




Visualizing Neural Networks



Understanding the Components, Structure, and Learning Process

Neural Network Components I

Layers and Neurons


Layers

- Organize computation into stages within the network.
- Input layer receives features from the dataset.
- Hidden layers transform inputs into useful internal patterns.
- Output layer produces the prediction or classification.

Neurons

- Basic processing units of a neural network.
- Receive inputs, combine them, and apply an activation function.
- Pass useful signals forward to the next layer.
- Many neurons together help detect complex relationships.

How layers and neurons form a simple network



Neural Network Components II

Weights and Activation Functions

Weights

- Represent the strength of a connection between neurons.
- Larger absolute weight means stronger influence.
- Updated during training to reduce prediction error.
- Bias terms shift a neuron's response.

Activation Functions

- Introduce non-linearity into the model.
- Help neural networks learn patterns beyond straight-line relationships.
- Common examples: ReLU, sigmoid, and tanh.

Neuron computation

$$z = w \cdot x + b$$

$$a = \text{activation}(z)$$

Weights scale inputs; activation decides what passes forward.


Learned weights

Positive and negative values amplify or reduce incoming signals.

Activation output

The neuron decides how strongly it should pass information forward.

Example activation behavior



Neural Network Components III

Loss Functions and Optimization Algorithms


Loss Functions

- Measure the difference between prediction and actual answer.
- Turn model error into a numeric value.
- Examples: mean squared error and cross-entropy.


Optimization Algorithms

- Use the loss value to update weights and biases.
- Move the model toward better parameter values step by step.
- Popular optimizers include gradient descent, Adam, and RMSprop.

Training loop: error measurement to improvement



Loss should decrease over time






Why this matters

Loss tells the model how wrong it is. The optimizer uses that signal to improve predictions on the next pass.

Overall Architecture and Data Flow

From input features to prediction and training feedback

This flowchart shows how information moves through a neural network during prediction and training.



Summary and Key Insights

Why visualizing neural networks matters

Importance of visualization

- Makes abstract AI concepts easier to understand and explain.
- Shows how components connect and how data flows from input to prediction.
- Helps both technical and nontechnical audiences follow the training process.

Key insights gained

- Layers and neurons create the structure that enables learning.
- Weights, activation, loss, and optimization each play a distinct role.
- Feature choice, noise, network size, and training settings affect model behavior.

Layers

Neurons

Weights

Activation

Loss

Optimization

References

Sources used for this artifact

- Goodfellow, I., Bengio, Y., & Courville, A. (2016). Deep learning. MIT Press. <https://www.deeplearningbook.org/>
- TensorFlow Playground. (n.d.). Neural network playground. <https://playground.tensorflow.org/>
- IBM. (n.d.). What are neural networks? <https://www.ibm.com/topics/neural-networks>
- Google Developers. (n.d.). Machine learning crash course. <https://developers.google.com/machine-learning/crash-course>